


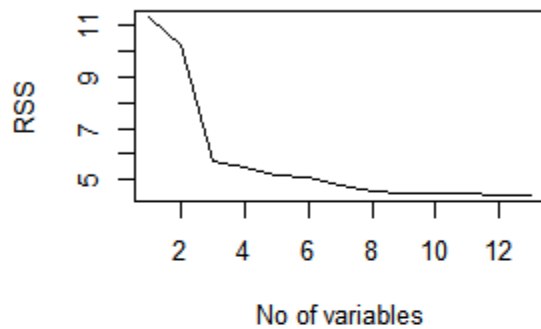


CS 6301.012  
Advanced Computing for Data Science  
Assignment VI

Sushmitha Mohan Raj sxm144630

Sreesha Nagaraj sxn146630

## LPGA Winnings:



```
> subsets_full = regsubsets(lpga2009$v14 ~., data=lpga2009)
> subsets_full.summary = summary(subsets_full)
> subsets_full.summary
Subset selection object
Call: regsubsets.formula(lpga2009$v14 ~ ., data = lpga2009)
13 variables (and intercept)
    Forced in Forced out
V1      FALSE      FALSE
V2      FALSE      FALSE
V3      FALSE      FALSE
V4      FALSE      FALSE
V5      FALSE      FALSE
V6      FALSE      FALSE
V7      FALSE      FALSE
V8      FALSE      FALSE
V9      FALSE      FALSE
V10     FALSE      FALSE
V11     FALSE      FALSE
V12     FALSE      FALSE
V13     FALSE      FALSE
1 subsets of each size up to 8
Selection Algorithm: exhaustive
      v1 v2 v3 v4 v5 v6 v7 v8 v9 v10 v11 v12 v13
```

```

1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " * " " " "
2 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " * " " " "
3 ( 1 ) " " " " " " " " " * " * " " " " " " " " " " " " " " * " " "
4 ( 1 ) " " " " " " " " " * " * " " " " " " " * " " " " " " " * " " "
5 ( 1 ) " " " " " " " " " * " * " " " " " " " * " " " " " * " * " " "
6 ( 1 ) " " " " " " " " " * " * " * " " " " " * " " " " " * " * " " "
7 ( 1 ) " " " * " * " * " * " " " " " " " * " " " " " * " * " " "
8 ( 1 ) " " " * " * " * " * " * " " " " " * " " " * " * " " "
> subsets_full_13 = regsubsets(lpga2009$v14 ~., data=lpga2009,nvmax = 14)
> subsets_full_13.summary = summary(subsets_full_13)
> subsets_full_13.summary
Subset selection object
Call: regsubsets.formula(lpga2009$v14 ~ ., data = lpga2009, nvmax = 14)
13 variables (and intercept)
      Forced in Forced out
V1          FALSE      FALSE
V2          FALSE      FALSE
V3          FALSE      FALSE
V4          FALSE      FALSE
V5          FALSE      FALSE
V6          FALSE      FALSE
V7          FALSE      FALSE
V8          FALSE      FALSE
V9          FALSE      FALSE
V10         FALSE      FALSE
V11         FALSE      FALSE
V12         FALSE      FALSE
V13         FALSE      FALSE
1 subsets of each size up to 13
Selection Algorithm: exhaustive
      V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13
1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " * " " "
2 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " * " " "
3 ( 1 ) " " " " " " " " " * " * " " " " " " " " " " " " " " * " " "
4 ( 1 ) " " " " " " " " " * " * " " " " " " " * " " " " " " * " " "
5 ( 1 ) " " " " " " " " " * " * " " " " " " " * " " " " * " " "
6 ( 1 ) " " " " " " " " " * " * " * " " " " " * " " " " * " " "
7 ( 1 ) " " " * " * " * " * " " " " " " " * " " " " " * " " "
8 ( 1 ) " " " * " * " * " * " * " " " " " * " " " " " * " * " "
9 ( 1 ) " * " * " * " * " * " * " " " " * " " " " " * " * " "
10 ( 1 ) " * " * " * " * " * " * " * " * " " " " " " * " * " "
11 ( 1 ) " * " * " * " * " * " * " * " * " " " " " " * " * " *
12 ( 1 ) " * " * " * " * " * " * " * " * " * " " " " * " * " *
13 ( 1 ) " * " * " * " * " * " * " * " * " * " * " * " * " *
> par(mfrow=c(2,2))
> plot(subsets_full_13$rss,xlab = "No of variables", ylab = "RSS", type = "l")
> names(subsets_full_13.summary)
[1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
> plot(subsets_full_13.summary$rss,xlab = "No of variables", ylab = "RSS", ty
pe = "l")
> plot(subsets_full_13.summary$rss,xlab = "No of variables", ylab = "RSS", ty
pe = "l")
> par(mfrow=c(2,2))
> plot(subsets_full_13.summary$rss,xlab = "No of variables", ylab = "RSS", ty
pe = "l")





```

```

> plot(subsets_full_13.summary$adjr2,xlab = "No of variables", ylab = "RSS",
type = "l")
> which.max(subsets_full_13.summary)
Error in which.max(subsets_full_13.summary) :
  (list) object cannot be coerced to type 'double'
> which.max(subsets_full_13.summary$adjr2)
[1] 9
> points(9,subsets_full_13.summary$adjr2[9],col="red",cex=2,pch=20)
> plot(subsets_full_13.summary$cp,xlab = "No of variables", ylab = "Cp", type
= "l")
> which.min(subsets_full_13.summary$cp)
[1] 8
> points(8,subsets_full_13.summary$cp[8],col="red",cex=2,pch=20)
> plot(subsets_full_13.summary$bic,xlab = "No of variables", ylab = "BIC", ty
pe = "l")
> which.min(subsets_full_13.summary$bic)
[1] 8
> points(8,subsets_full_13.summary$bic[8],col="red",cex=2,pch=20)

> plot(subsets_full_13,scale="r2")
> plot(subsets_full_13,scale="adjr2")
> plot(subsets_full_13,scale="Cp")
> plot(subsets_full_13,scale="bic")

```







### Forward Subset Selection:





```
> subset_fwd = regsubsets(V14~., data=lpga2009, method = "forward",nvmax = 13
)
> subset_fwd.summary = summary(subset_fwd)
> subset_fwd.summary
Subset selection object
Call: regsubsets.formula(V14 ~ ., data = lpga2009, method = "forward",
  nvmax = 13)
13 Variables (and intercept)
   Forced in Forced out
V1      FALSE      FALSE
V2      FALSE      FALSE
V3      FALSE      FALSE
V4      FALSE      FALSE
V5      FALSE      FALSE
V6      FALSE      FALSE
V7      FALSE      FALSE
V8      FALSE      FALSE
V9      FALSE      FALSE
V10     FALSE      FALSE
V11     FALSE      FALSE
V12     FALSE      FALSE
V13     FALSE      FALSE
1 subsets of each size up to 13
Selection Algorithm: forward
   V1  V2  V3  V4  V5  V6  V7  V8  V9  V10 V11 V12 V13
1 ( 1 ) " " " " " " " " " " " " " " " " "*" " "
2 ( 1 ) " " " " " " " " " " " " " " " " "*" " "
3 ( 1 ) " " " " " " "*" " " " " " " " " " " "*" " "
4 ( 1 ) " " " " " " "*" "*" " " " " " " " " "*" " "
5 ( 1 ) " " " " " " "*" "*" "*" " " " " " " " "*" " "
6 ( 1 ) " " " " " " "*" "*" "*" " " "*" " " " " "*" " "
7 ( 1 ) " " " " " " "*" "*" "*" " " "*" " " " " "*" " "
8 ( 1 ) " " "*" " " " "*" "*" "*" " " "*" " " " " "*" " "
9 ( 1 ) " " "*" "*" "*" "*" "*" " " "*" " " " " "*" " "
10 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" " " " " "*" " "
11 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" " " "*" " " " " "*" " "
12 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" " " "*" " " " " "*" " "
13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" " " "*" " " " " "*" " "

> par(mfrow=c(2,2))
> plot(subset_fwd.summary$rss,xlab = "No of variables", ylab = "RSS", type =
"l")
> plot(subset_fwd.summary$adjr2,xlab = "No of variables", ylab = "AdjR2", typ
e = "l")
> which.max(subset_fwd.summary$adjr2)
[1] 9
> points(9, subset_fwd.summary$adjr2[9], col="red", cex=2,pch=20)
> plot(subset_fwd.summary$cp,xlab = "No of variables", ylab = "Cp", type = "l
")
> which.min(subset_fwd.summary$cp)
[1] 9
> points(9, subset_fwd.summary$cp[9], col="red", cex=2,pch=20)
> plot(subset_fwd.summary$bic,xlab = "No of variables", ylab = "bic", type =
"l")
```

```
> which.min(subset_fwd.summary$bic)
[1] 9
> points(9, subset_fwd.summary$bic[9], col="red", cex=2, pch=20)
```



```
> plot(subset_fwd, scale = "r2")
> plot(subset_fwd, scale = "adjr2")
> plot(subset_fwd, scale = "Cp")
> plot(subset_fwd, scale = "bic")
```



## Backward Stepwise Selection

```
subset_bwd = regsubsets(V14~., data=lpga2009, method = "backward", nvmax = 13)
> subset_bwd.summary = summary(subset_bwd)
> subset_bwd.summary
subset selection object
Call: regsubsets.formula(V14 ~ ., data = lpga2009, method = "backward",
  nvmax = 13)
```

13 Variables (and intercept)

	Forced in	Forced out
V1	FALSE	FALSE
V2	FALSE	FALSE
V3	FALSE	FALSE
V4	FALSE	FALSE
V5	FALSE	FALSE
V6	FALSE	FALSE
V7	FALSE	FALSE
V8	FALSE	FALSE
V9	FALSE	FALSE
V10	FALSE	FALSE
V11	FALSE	FALSE
V12	FALSE	FALSE
V13	FALSE	FALSE

1 subsets of each size up to 13

Selection Algorithm: backward





		v1	v2	v3	v4	v5	v6	v7	v8	v9	v10	v11	v12	v13
1	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
2	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
3	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
4	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
5	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
6	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
7	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
8	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
9	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
10	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
11	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
12	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
13	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "

```





points(9, subset_bwd.summary$adjr2[9], col="red", cex=2,pch=20)
> plot(subset_bwd.summary$rss,xlab = "No of variables", ylab = "RSS", type =
"l")
> plot(subset_bwd.summary$adjr2,xlab = "No of variables", ylab = "adjr2", typ
e = "l")
> which.max(subset_bwd.summary$adjr2)
[1] 9
> points(9, subset_bwd.summary$adjr2[9], col="red", cex=2,pch=20)
> plot(subset_bwd.summary$cp,xlab = "No of variables", ylab = "Cp", type = "l
")
> which.min(subset_bwd.summary$cp)
[1] 8
> points(8, subset_bwd.summary$cp[8], col="red", cex=2,pch=20)
> plot(subset_bwd.summary$bic,xlab = "No of variables", ylab = "bic", type =
"l")
> which.min(subset_bwd.summary$bic)
[1] 8
> points(8, subset_bwd.summary$bic[8], col="red", cex=2,pch=20)

```





```
> plot(subset_bwd, scale = "r2")  
> plot(subset_bwd, scale = "adjr2")  
> plot(subset_bwd, scale = "Cp")  
> plot(subset_bwd, scale = "bic")
```



```
> coef(subsets_full_13,13)
(Intercept)          v1          v2          v3          v4
v5
 7.130639e+01 -2.839141e-04 -1.017259e-02 -1.272097e-02 -9.737279e-02  4.5076
19e-01
          v6          v7          v8          v9          v10
v11
-5.523654e-03 -6.760145e-08 -1.271894e-01 -1.042326e-02  5.079487e-01  9.9188
72e-03
          v12          v13
-3.400786e-02  6.022010e-03
> coef(subset_fwd,13)
(Intercept)          v1          v2          v3          v4
v5
 7.130639e+01 -2.839141e-04 -1.017259e-02 -1.272097e-02 -9.737279e-02  4.5076
19e-01
          v6          v7          v8          v9          v10
v11
-5.523654e-03 -6.760145e-08 -1.271894e-01 -1.042326e-02  5.079487e-01  9.9188
72e-03
          v12          v13
-3.400786e-02  6.022010e-03
> coef(subset_bwd,13)
```

(Intercept)	V1	V2	V3	V4
V5 7.130639e+01 19e-01	-2.839141e-04	-1.017259e-02	-1.272097e-02	-9.737279e-02 4.5076
V6	V7	V8	V9	V10
V11 -5.523654e-03 72e-03	-6.760145e-08	-1.271894e-01	-1.042326e-02	5.079487e-01 9.9188
V12	V13			
-3.400786e-02	6.022010e-03			

### Using Cross Validation Approach

```
> set.seed(1)
> train = sample(c(TRUE,FALSE), nrow(lpga2009),rep=TRUE)
> test = (!train)
> subset_train_full = regsubsets(V14 ~., data=lpga2009[train,],nvmax = 13)
> train_full_Summary = summary(subset_train_full)
> train_full_Summary
```

Subset selection object

Call: regsubsets.formula(V14 ~ ., data = lpga2009[train, ], nvmax = 13)

13 variables (and intercept)

Forced in Forced out

V1	FALSE	FALSE
V2	FALSE	FALSE
V3	FALSE	FALSE
V4	FALSE	FALSE
V5	FALSE	FALSE
V6	FALSE	FALSE
V7	FALSE	FALSE
V8	FALSE	FALSE
V9	FALSE	FALSE
V10	FALSE	FALSE
V11	FALSE	FALSE
V12	FALSE	FALSE
V13	FALSE	FALSE





1 subsets of each size up to 13

Selection Algorithm: exhaustive

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13
1 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
2 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
3 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
4 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
5 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
6 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
7 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
8 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
9 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
10 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
11 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
12 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"
13 ( 1 )	"	"	"	"	"	"	"	"	"	"	"	"	"

```
> plot(subset_train_full, scale="r2")
> plot(subset_train_full, scale="adjr2")
```

```
> plot(subset_train_full, scale="Cp")
> plot(subset_train_full, scale="bic")
```



```
> test.mat=model.matrix(V14 ~., data= lpga2009[test,])
> val.errors =rep(NA ,13)
> for(i in 1:13){
+   coef= coef(subset_train_full,id=i)
+   pred = test.mat[,names(coef)]%*%coef
+   val.errors[i]=mean((lpga2009$V14[test]-pred)^2)
+ }
> val.errors
[1] 0.08415571 0.10273877 0.04445451 0.04549853 0.05412799 0.05313093 0.0551
6970
[8] 0.05017166 0.05133944 0.05329075 0.05279580 0.05452616 0.05458574
> which.min(val.errors)
[1] 3
```

We find that the best model is the one that contains three variables.

```
> coef(subset_train_full,3)
(Intercept)          v4          v5          v12
66.67275719 -0.12176710  0.53265360 -0.03739409
```

**Coefficients on the full data set**

```
> subset_full = regsubsets(V14 ~., data=lpga2009,nvmax = 13)
```


```
> coef(subset_full,3)
(Intercept)          V4          V5          V12
66.88662715 -0.11665632  0.51442503 -0.03802276
```

### Cross-Validation

```
k=10
> set.seed(1)
> folds=sample (1:k,nrow(lpga2009),replace =TRUE)
> cv.errors =matrix (NA ,k,13, dimnames =list(NULL , paste (1:13) ))
> for(j in 1:k){
+   best.fit =regsubsets(V14~.,data=lpga2009[folds!=j,], nvmax=13)
+   for(i in 1:13){
+     pred=predict.regsubsets(best.fit,lpga2009[folds==j,],id=i)
+     cv.errors [j,i]=mean((V14[folds==j]-pred)^2)
+   }
+ }
>
> mean.cv.errors =apply(cv.errors ,2, mean)
> mean.cv.errors
```

	1	2	3	4	5	6	7
	0.07574207	0.08217386	0.04109246	0.04235880	0.04366251	0.04562376	0.04044721
	8	9	10	11	12	13	
	0.03697168	0.03924182	0.04003496	0.04025059	0.03999827	0.03983661	

## The Cross-Validation select model with 8 variables



```
> subset.best = regsubsets(V14~.,data=lpga2009,nvmax = 13)
> coef(subset.best,8)
      (Intercept)          v2          v3          v4          v5          v6
71.993534526 -0.010499992 -0.012789745 -0.096948026  0.464373516 -0.005613287
          v8          v11          v12
-0.129541406  0.019951333 -0.035587780
```

## Ridge Regression:

```
> grid =10^ seq (10,-2, length =100)
> ridge.mod =glmnet (x,y,alpha =0, lambda =grid)
> dim(coef(ridge.mod ))
[1] 14 100
> ridge.mod$lambda [50]
[1] 11497.57
> coef(ridge.mod)[,50]
      (Intercept)          v1          v2          v3          v4          v5
7.267457e+01 -5.163723e-07 -6.036248e-06 -4.671971e-06 -2.094081e-05  7.2719
69e-05
```

```

          v6          v7          v8          v9          v10
v11
-3.193423e-06 -2.445436e-10 -7.333001e-05 -1.610827e-05  2.166955e-03 -1.5160
85e-05
          v12          v13
-7.247062e-06 -4.785361e-06
> sqrt(sum(coef(ridge.mod)[ -1 ,50]^2) )
[1] 0.002169662

```

```

> sqrt(sum(coef(ridge.mod)[ -1 ,50]^2) )
[1] 0.002169662
> ridge.mod$lambda [60]
[1] 705.4802
> coef(ridge.mod)[,60]
(Intercept)          v1          v2          v3          v4
v5
 7.266415e+01 -8.330121e-06 -9.740309e-05 -7.567890e-05 -3.382856e-04  1.1742
84e-03
          v6          v7          v8          v9          v10
v11
-5.148019e-05 -3.945858e-09 -1.183264e-03 -2.592571e-04  3.498474e-02 -2.4450
67e-04
          v12          v13
-1.170228e-04 -7.713420e-05
> sqrt(sum(coef(ridge.mod)[ -1 ,60]^2) )
[1] 0.03502842

```

```

> predict(ridge.mod ,s=50, type ="coefficients")[1:14 ,]
(Intercept)          v1          v2          v3          v4
v5
 7.253018e+01 -1.045822e-04 -1.227944e-03 -1.004749e-03 -4.330192e-03  1.4955
21e-02
          v6          v7          v8          v9          v10
v11
-6.402104e-04 -4.968126e-08 -1.490318e-02 -3.146094e-03  4.437079e-01 -3.0541
22e-03
          v12          v13
-1.487303e-03 -9.553873e-04

```

```

set.seed(1)
> train=sample (1: nrow(x), nrow(x)/2)
> test=(- train )
> y.test=y[test]
ridge.mod =glmnet (x[train ,],y[train],alpha =0, lambda =grid ,
+                 thresh =1e-12)
ridge.pred=predict (ridge.mod ,s=4, newx=x[test ,])
> mean(( ridge.pred -y.test)^2)
[1] 0.1745175

```

```

MSE is 0.1745
mean(( mean(y[train ]) -y.test)^2)
[1] 1.086542

```

```

> ridge.pred=predict (ridge.mod ,s=1e10 ,newx=x[test ,])
> mean(( ridge.pred -y.test)^2)

```

```
[1] 1.086542
```

```
> ridge.pred=predict (ridge.mod ,s=0, newx=x[test ,], exact=T)
```

```
> mean(( ridge.pred -y.test)^2)
```

```
[1] 0.05601155
```

```
> lm(y~x, subset =train)
```

```
Call:
```

```
lm(formula = y ~ x, subset = train)
```

```
Coefficients:
```

(Intercept)	xv1	xv2	xv3	xv4	xv5
7.506e+01	3.913e-04	-6.025e-03	-1.111e-02	-8.117e-02	4.008e-01
xv6	xv7	xv8	xv9	xv10	xv11
7.283e-04	-4.746e-08	-2.245e-01	-7.642e-02	-1.106e+00	1.256e-03
xv12	xv13				
-4.611e-02	2.861e-02				

```
> predict (ridge.mod ,s=0, exact =T,type="coefficients") [1:14 ,]
```


(Intercept)	v1	v2	v3	v4	
v5					
7.505955e+01	3.913258e-04	-6.024746e-03	-1.111252e-02	-8.117081e-02	4.0074
47e-01					
v6	v7	v8	v9	v10	
v11					
7.283714e-04	-4.745528e-08	-2.245362e-01	-7.642078e-02	-1.105755e+00	1.2554
22e-03					
v12	v13				
-4.610965e-02	2.861557e-02				

```
> set.seed (1)
```

```
> cv.out =cv.glmnet (x[train ,],y[train],alpha =0)
```

```
> plot(cv.out)
```





```
> bestlam =cv.out$lambda.min
> bestlam
[1] 0.1317802
```


```
ridge.pred=predict (ridge.mod ,s=bestlam ,newx=x[test ,])
> mean(( ridge.pred -y.test)^2)
[1] 0.04282599
```

**There is an improvement in the MSE**


```
> out=glmnet (x,y,alpha =0)
> predict (out ,type="coefficients",s=bestlam )[1:14 ,]
      (Intercept)          v1          v2          v3          v4
v5
 7.016981e+01 -2.894303e-04 -1.042156e-02 -1.622324e-02 -7.532142e-02  2.6534
45e-01
          v6          v7          v8          v9          v10
v11
-5.848681e-03 -2.211937e-07 -1.243824e-01  1.620534e-02  3.278224e+00 -9.3285
98e-03
          v12          v13
-2.114120e-02 -1.256978e-03
```

### The Lasso Regression

```
> lasso.mod = glmnet (x[train ,],y[train],alpha =1, lambda =grid)
> plot(lasso.mod)
```




```
> set.seed (1)
> cv.out =cv.glmnet (x[train ,],y[train],alpha =1)
> plot(cv.out)
```



```
> bestlam =cv.out$lambda.min
> bestlam
[1] 0.005976074

> lasso.pred=predict (lasso.mod ,s=bestlam ,newx=x[test ,])
> mean(( lasso.pred -y.test)^2)
[1] 0.04374007
```

This MSE is .001 more than ridge



```

out=glmnet (x,y,alpha =1, lambda =grid)
> lasso.coef=predict(out ,type ="coefficients",s=bestlam )[1:14 ,]
> lasso.coef
  (Intercept)      v1      v2      v3      v4
v5
  6.948794e+01 -3.036679e-05 -6.106838e-03 -8.053118e-03 -9.695680e-02  3.9741
24e-01
      v6      v7      v8      v9      v10
v11
 -4.303790e-03 -8.161627e-08 -7.266772e-02  4.534686e-03  1.278453e+00  0.0000
00e+00
      v12      v13
-3.164070e-02  0.000000e+00

```

## Communities and Crime:

The data set contains 122 predictors. The response variable is "ViolentCrimesPerPop". The cells with "?" were removed using the following code.

```
crimeData=read.csv("C:\\Users\\Sushmitha\\Documents\\Third Semester\\Comp Mthds Data  
Science\\Assignment VI\\Assign_6-Data\\Assign_6-Data\\communitiesCrimeDataset.csv",",", header=  
TRUE)
```

```
idx <- crimeData == "?"
```

```
is.na(crimeData) <- idx
```

```
crimeData=na.omit(crimeData)
```

The final data set had 319 rows.

The Best selection model could not be executed due to higher count of predictors. So following are the models built using the "forward" and "backward" models.

### #Forward Selection

```
regfit.fwd=regsubsets(ViolentCrimesPerPop~, crimeData, nvmax = 55, method = "forward")
```

```
head(summary(regfit.fwd))
```

\$rsq

```
[1] 0.5759351 0.6251904 0.6521741 0.6665235 0.6804758 0.6924569 0.7037983 0.7146353  
[9] 0.7272709 0.7378819 0.7460173 0.7532894 0.7596104 0.7651628 0.7705080 0.7758095  
[17] 0.7811379 0.7861956 0.7913152 0.7964764 0.8018642 0.8062093 0.8113085 0.8171155  
[25] 0.8226798 0.8269001 0.8314047 0.8350404 0.8382924 0.8414197 0.8444722 0.8474111  
[33] 0.8502837 0.8533304 0.8562424 0.8592216 0.8622088 0.8646632 0.8674861 0.8700544  
[41] 0.8724990 0.8747540 0.8770173 0.8790561 0.8811733 0.8831980 0.8851595 0.8871262  
[49] 0.8891868 0.8915737 0.8938489 0.8961034 0.8981283 0.9002118 0.9024843 0.9044487
```

\$rss

```
[1] 10.298648 9.102456 8.447142 8.098659 7.759820 7.468852 7.193421 6.930239  
[9] 6.623375 6.365680 6.168108 5.991503 5.837993 5.703149 5.573338 5.444588  
[17] 5.315186 5.192356 5.068023 4.942681 4.811837 4.706314 4.582477 4.441450  
[25] 4.306319 4.203826 4.094428 4.006135 3.927158 3.851210 3.777078 3.705704
```

[33] 3.635943 3.561952 3.491232 3.418880 3.346335 3.286729 3.218172 3.155800  
[41] 3.096433 3.041668 2.986702 2.937190 2.885772 2.836600 2.788965 2.741203  
[49] 2.691158 2.633193 2.577938 2.523186 2.474010 2.423412 2.368221 2.320516

#### \$adjr2

[1] 0.5745974 0.6228182 0.6488615 0.6622754 0.6753716 0.6865426 0.6971314 0.7072710  
[9] 0.7193274 0.7293716 0.7369170 0.7436144 0.7493643 0.7543479 0.7591471 0.7639319  
[17] 0.7687769 0.7733674 0.7780543 0.7828171 0.7878546 0.7918059 0.7965969 0.8021862  
[25] 0.8075500 0.8114871 0.8157619 0.8191132 0.8220656 0.8249009 0.8276730 0.8303382  
[33] 0.8329481 0.8357713 0.8384632 0.8412499 0.8440655 0.8462960 0.8489627 0.8513572  
[41] 0.8536270 0.8556948 0.8577873 0.8596344 0.8615864 0.8634447 0.8652425 0.8670597  
[49] 0.8690015 0.8713449 0.8735728 0.8757928 0.8777539 0.8798005 0.8820913 0.8840255

#### \$cp

[1] -14708.212 -13034.435 -12116.579 -11627.546 -11151.989 -10743.338 -10356.400  
[8] -9986.582 -9555.714 -9193.565 -8915.441 -8666.621 -8450.078 -8259.623  
[15] -8076.200 -7894.262 -7711.412 -7537.747 -7361.982 -7184.806 -6999.940  
[22] -6850.463 -6675.391 -6476.294 -6285.437 -6140.194 -5985.302 -5859.904  
[29] -5747.529 -5639.385 -5533.779 -5432.028 -5332.531 -5227.123 -5126.285  
[36] -5023.167 -4919.779 -4834.476 -4736.662 -4647.492 -4562.521 -4483.983  
[43] -4405.163 -4333.965 -4260.106 -4189.383 -4120.809 -4052.058 -3980.116  
[50] -3897.105 -3817.881 -3739.360 -3668.634 -3595.918 -3516.785 -3448.113

#### \$bic


[1] -262.1298 -295.7510 -313.8202 -321.4944 -329.3631 -335.7894 -342.0104 -348.1352  
[9] -356.8173 -363.7112 -368.0038 -371.5055 -374.0200 -375.7094 -377.2890 -378.9795  
[17] -380.8875 -382.5807 -384.5470 -386.7705 -389.5638 -390.8721 -393.6131 -397.8195  
[25] -401.9106 -403.8296 -406.4758 -407.6649 -408.2512 -408.7156 -409.1507 -409.4713  
[33] -409.7686 -410.5620 -411.1941 -412.1093 -413.1858 -413.1539 -414.1130 -414.5911

[41] -414.8842 -414.8114 -414.8636 -414.4310 -414.2995 -414.0168 -413.6541 -413.3992


[49] -413.5116 -414.6926 -415.6925 -416.7755 -417.2887 -418.1154 -419.6991 -420.4254

The file `output_FwdSelection.txt` attached with the assignment shows the summary of the forward model.


`plot(regfit.sum_fwd$rss, xlab = "No of variables", ylab = "RSS", type = "l")`



```
points(56,regfit.sum_fwd$rss[56], col = "red", cex=2, pch=20)
```



```
plot(regfit.sum_fwd$adjr2, xlab = "No of variables", ylab = "ADJR2", type = "l")
```





From the above plots we can infer that the model with nvmax=55, the adjusted R square with one predictor gives the best model. And the RSS and adjusted R square are inversely proportional.

### #Backward Selection

```
regfit.bwd=regsubsets(ViolentCrimesPerPop~., crimeData, nvmax = 55, method = "backward")
sink(file="output_BackwardSelection.txt")
regfit.sum_bwd=summary(regfit.bwd)
head(summary(regfit.bwd))
```

\$rsq

```
[1] 0.4885507 0.5675394 0.6077228 0.6344460 0.6416150 0.6493676 0.6575348 0.6681325
[9] 0.6772358 0.6849204 0.6898604 0.6903921 0.6976386 0.7057232 0.7067725 0.7140608
[17] 0.7191327 0.7242183 0.7298332 0.7352423 0.7403849 0.7453741 0.7490500 0.7527328
[25] 0.7531287 0.7555004 0.7595096 0.7639418 0.7682604 0.7722317 0.7746182 0.7780789
[33] 0.7799249 0.7830703 0.7835444 0.7847145 0.7872475 0.7899424 0.7913577 0.7941932
[41] 0.7968644 0.7992922 0.8010438 0.8031356 0.8046045 0.8065973 0.8089763 0.8118458
[49] 0.8144644 0.8166573 0.8185472 0.8207094 0.8225997 0.8243216 0.8262531 0.8269956
```

\$rss

```
[1] 12.420826 10.502543 9.526666 8.877679 8.703576 8.515300 8.316954 8.059584
[9] 7.838505 7.651881 7.531911 7.518997 7.343012 7.146674 7.121190 6.944190
[17] 6.821016 6.697508 6.561149 6.429785 6.304894 6.183728 6.094458 6.005019
[25] 5.995406 5.937807 5.840440 5.732803 5.627922 5.531479 5.473522 5.389475
[33] 5.344644 5.268257 5.256742 5.228326 5.166810 5.101364 5.066992 4.998131
[41] 4.933260 4.874299 4.831759 4.780959 4.745287 4.696890 4.639115 4.569427
[49] 4.505834 4.452579 4.406682 4.354170 4.308263 4.266447 4.219538 4.201507
```

\$adjr2

[1] 0.4869373 0.5648023 0.6039868 0.6297893 0.6358900 0.6426246 0.6498266 0.6595682  
[9] 0.6678349 0.6746905 0.6787479 0.6782506 0.6847511 0.6921710 0.6922563 0.6989117  
[17] 0.7032698 0.7076714 0.7126654 0.7174734 0.7220283 0.7264492 0.7294844 0.7325478  
[25] 0.7320646 0.7337299 0.7371961 0.7411500 0.7450063 0.7485058 0.7502738 0.7532486  
[33] 0.7544426 0.7570998 0.7567743 0.7572312 0.7592339 0.7614346 0.7621927 0.7645807  
[41] 0.7667973 0.7687497 0.7699343 0.7715224 0.7723965 0.7738895 0.7758467 0.7783962  
[49] 0.7806679 0.7824515 0.7838876 0.7856601 0.7871197 0.7883874 0.7899182 0.7900176

\$cp


[1] -17674.131 -14991.173 -13625.303 -12716.290 -12470.966 -12205.836 -11926.631  
[8] -11564.934 -11253.959 -10991.136 -10821.468 -10801.420 -10553.466 -10277.067  
[15] -10239.452 -9990.080 -9815.933 -9641.322 -9448.748 -9263.156 -9086.610  
[22] -8915.271 -8788.509 -8661.510 -8646.075 -8563.576 -8425.498 -8273.065  
[29] -8124.487 -7987.699 -7904.699 -7785.236 -7720.581 -7611.824 -7593.731  
[36] -7552.018 -7464.044 -7370.578 -7320.541 -7222.302 -7129.639 -7045.236  
[43] -6983.783 -6910.785 -6858.930 -6789.292 -6706.547 -6607.152 -6516.275  
[50] -6439.847 -6373.702 -6298.313 -6232.153 -6171.712 -6104.153 -6076.953

\$bic


[1] -202.3613 -250.1106 -275.4552 -292.1969 -292.7499 -293.9610 -295.7142 -299.9765  
[9] -303.0839 -305.0055 -304.2814 -299.0636 -300.8535 -303.7339 -299.1082 -301.3721  
[17] -301.3161 -301.3799 -302.1765 -302.8630 -303.3549 -303.7799 -302.6534 -301.6044  
[25] -296.3503 -293.6646 -293.1736 -293.3424 -293.4673 -293.2160 -290.8109 -289.9820  
[33] -286.8814 -285.7083 -280.6412 -276.6050 -274.6154 -272.9167 -269.3081 -267.9079  
[41] -266.3101 -264.3805 -261.4116 -259.0180 -255.6420 -253.1469 -251.3299 -250.3931  
[49] -249.0987 -247.1262 -244.6663 -242.7253 -240.3413 -237.6874 -235.4490 -231.0499

The file output\_ForwardSelection.txt attached with the assignment shows the summary of the forward model.


```
plot(regfit.sum_bwd$rss, xlab = "No of variables", ylab = "RSS", type = "l")
```



```
points(56,regfit.sum_bwd$rss[56], col = "red", cex=2, pch=20)
```



```
plot(regfit.sum_bwd$adjr2, xlab = "No of variables", ylab = "ADJR2", type = "l")
```



From the above plot we can infer that the model with  $nvmax=55$ , the adjusted R square with one predictor gives the best model. And the RSS and adjusted R square are inversely proportional.